



HMIS — Hiring Manager Interview System System Design Document

AI-augmented role discovery, hiring-pack generation, and candidate evaluation

Prepared for	ZURU AI Integration Internship Task
Prepared by	Aditya R Aryamaan
Demonstrated stack	Telegram, n8n, OpenRouter, Gotenberg

Contents

1	Executive summary	1
2	Scope and design principles	1
2.1	In-scope proof of concept	1
2.2	Architecture principles	1
3	System architecture	2
3.1	Logical component view	2
3.2	Component responsibilities	3
3.3	Canonical information contract	3
4	Information Extraction Framework	3
4.1	Nine key hiring-manager questions	3
4.2	Coverage and completion rules	4
4.3	Hiring Brief Readiness calculation	4
4.4	Handling vague responses	5
5	Capturing technical skills and cultural fit	5
5.1	ZURU DNA translation examples	6
6	Risk Assessment Matrix	7
6.1	Recommended human checkpoints	7
7	Process Flow Diagrams	8
7.1	Current hiring workflow and pain points	8
7.2	Proposed AI-augmented workflow	8
8	Candidate evaluation logic	9
8.1	Score, confidence, and review priority are separate	9
8.2	Post-interview reviewer actions	9
9	Operational controls, rollout, and metrics	9
9.1	Failed approaches and lessons learned	9
9.2	Security and governance controls	9
9.3	Phased rollout	10
9.4	POC validation evidence and pilot test plan	10
10	Traceability to assignment deliverables	11

1 Executive summary

The proposed Hiring Manager Interview System (HMIS) replaces an inconsistent sequence of meetings, document revisions, and disconnected screening judgments with one traceable requirements-to-evidence workflow. A hiring manager uses a familiar chat channel to define the role. The system checks nine required information areas, converts vague language into observable requirements, applies employment-type rules, and generates a job description, a hiring guide, and a machine-readable rubric. It then reports a deterministic Hiring Brief Readiness score with the strongest and least-defined areas. Talent Acquisition (TA) approves that rubric before it can be used for candidate review.

Candidate evaluation is asynchronous and takes place after an interview or structured screening response. The evaluator compares supplied evidence with the approved rubric criterion by criterion. It reports the matching anchors, exact supporting excerpts, an evidence-confidence category, and a deterministic human-review priority. It does not issue an overall score, ranking, or hire/reject decision.

Assignment baseline	Solution
2–3 meetings and 1–2 hours to define a role	Nine-area adaptive interview with a manager-confirmed review brief
40% of job descriptions require major revision after posting	Requirements, priorities, and unresolved decisions are confirmed before document generation
Five minutes per initial application review	Reusable screening questions and an evidence-anchored hiring rubric support consistent review
25% pass screening but fail basic requirements during the first interview	Every essential criterion is connected to a screening question and observable evidence anchors
Inconsistent evaluation against ZURU DNA	Relevant DNA themes are translated into observable, role-specific working behaviours
Must-have and nice-to-have requirements emerge late	Capabilities are explicitly classified as essential, preferred, or trainable during role discovery

Core design choice. The job description, hiring guide, and evaluator all derive from the same versioned role rubric. This prevents the candidate evaluator from silently inventing criteria that were never approved or advertised.

2 Scope and design principles

2.1 In-scope proof of concept

- Telegram inline-button entry point for nontechnical users, with free-text adaptive conversation.
- Local n8n orchestration and temporary conversation memory.
- OpenRouter-hosted language model for structured extraction and criterion-level evidence interpretation.
- Deterministic schema, coverage, employment-type, brief-readiness, approval, exact-quote, score, and confidence-ceiling validation.
- Local Gutenberg conversion to branded PDF outputs.
- Separate evaluator form accepting synthetic pasted evidence or TXT/VTT transcripts.

2.2 Architecture principles

1. **Ask for evidence, not adjectives.** Convert “great communicator” and “gets our vibe” into observable actions and role-relevant examples.
2. **Separate essentials, preferences, and trainable skills.** Each essential requires a reason tied to an outcome.
3. **One approved contract.** The versioned rubric connects discovery, questions, scoring anchors, and evaluation.

4. **Missing evidence is not failure.** It produces an unscored criterion and a reviewer action.
5. **AI proposes; humans own decisions.** Approval gates exist before publication, evaluation, and progression.
6. **Low-code and channel-independent.** n8n owns the workflow; Telegram is a replaceable adapter.
7. **Make programme rules explicit.** Employment type controls contract wording: internships use the approved summer programme pattern, permanent roles omit employment dates, and other fixed-term roles require an end condition.

3 System architecture

3.1 Logical component view




Figure 1: Logical component view based on the prototype concept. HMIS is the parent system; generation and candidate evaluation are its two modules.

3.2 Component responsibilities

Component	Responsibility	Control boundary
Telegram adapter	Menu buttons, callback routing, conversational messages, document delivery	Allowlisted user; no candidate transcript retained in chat
n8n orchestrator	State, branching, model calls, validation, document pipeline	Deterministic commands route before AI; credentials remain in n8n
Hiring-manager agent	Adaptive discovery, employment-type classification, and structured hiring-pack proposal	Cannot generate until review brief is shown and explicit confirmation is received
Schema validator	Enforces required fields, employment rules, 4–8 criteria, unique IDs, 5–7 questions, and criterion coverage	Rejects malformed, incomplete, or contract-inconsistent model output
Readiness calculator	Applies fixed weights to the nine discovery-area statuses and lists strong and under-defined areas	Code calculates the result; the model cannot invent or adjust the numeric score
Role rubric	Versioned source for outcomes, priorities, anchors, questions, flags, and checkpoints	TA approval required before candidate evaluation
Candidate evaluator	Maps supplied evidence to the approved anchors	Exact evidence excerpts required; no overall verdict is retained
Renderer/Gutenberg	Produces branded, portable PDFs	No external asset URL; logo is bundled locally
Human reviewers	Correct requirements, approve publication, inspect evidence, decide progression	Decision ownership never transfers to the model

3.3 Canonical information contract

The internal contract stores role facts separately from assessment criteria. Each criterion contains a priority, rationale, four observable anchors, green flags, and concerns. Screening questions reference criterion IDs. Readiness metadata records how clearly each discovery area was defined. The evaluator receives a snapshot of this same contract so requirement changes cannot be applied retroactively without a new version.

```
role: title, team, level, location, arrangement, employment_type, contract,
      outcomes, trainable, open_questions
criterion: id, name, essential|preferred, outcome_reason, DNA behaviour,
           anchors[1..4], green_flags[], concerns[]
question: id, criterion_ids[], text
readiness: score, band, area_statuses, strong_areas, needs_attention
metadata: rubric_id, version, generated_at, approval_status
```

4 Information Extraction Framework

4.1 Nine key hiring-manager questions

These are nine fixed *coverage areas*, not a rigid nine-turn script. Information already present in the manager's opening message is reused. The assistant asks one focused question per turn and follows up only where an answer is vague, contradictory, overly broad, or material.

#	Primary question	Captured fields	Adaptive handling
1	What role are we hiring, where does it sit, who does it report to, and why is it needed now?	Title, team, seniority, reporting line, business reason, hiring urgency	Resolve conflicting titles or stakeholders; do not silently choose.

#	Primary question	Captured fields	Adaptive handling
2	What are the three most important outcomes in the first 3–6 months, and how will we know they were achieved?	Prioritised outcomes, measures, time horizon	Converts “superstar” or “everything” into outcomes and trade-offs.
3	What will this person own, support, decide, and explicitly not own?	Responsibilities, boundaries, authority, day-one independence	Calibrates scope to intern, individual contributor, manager, or executive level.
4	Which capabilities must be present on day one, which are preferred, and which can be taught? Why is each essential?	Essential, preferred, trainable, rationale	Groups long tool lists; preserves OR requirements; challenges unsupported gates.
5	What role-specific work, tools, systems, audiences, or production constraints must they handle?	Technical or creative domain detail	Technical branch asks about systems, integrations, testing, permissions, and data; creative branch asks about audience, briefs, iteration, portfolio contribution, and production.
6	Who will they work with, and what supervision, review, onboarding, or training is available?	Stakeholders, collaboration, review cadence, support	Prevents entry-level roles from inheriting senior-level independence by accident.
7	What evidence should count, including equivalent or nontraditional experience?	Acceptable examples, work samples, projects, volunteering, equivalencies	Avoids prestige-only evidence; calibrates evidence to seniority.
8	Which observable working behaviours matter, and what would be a green flag or a concern to clarify?	Approved ZURU DNA behaviour, positive evidence, review concerns	Converts “vibe” into role-relevant behaviour; rejects personality and demographic inference.
9	What are the employment logistics and unresolved decisions?	Employment type, country, location, arrangement, contract, approved benefits, legal constraints, owners	A summer internship is normalised to 12 weeks commencing mid-November; permanent full-time documents omit employment dates; other fixed-term roles require a duration or end condition.

4.2 Coverage and completion rules

1. Parse the opening brief into the nine areas and mark each **confirmed** , **partial** , **conflicting** , or **unknown** .
2. Ask the highest-impact missing question, using role-specific wording.
3. For vague answers, request an example, expected action, context, and evidence of success.
4. For conflicts, preserve both positions and record the person who can resolve them.
5. Before generation, show a review brief containing facts, priorities, trainable skills, behaviours, constraints, and unresolved items.
6. Block generation if title, outcomes/scope, essentials, or location/arrangement remain unresolved.
7. After generation, calculate readiness from the nine recorded statuses and show the score, readiness band, strongest areas, and clarification areas with the document pack.

4.3 Hiring Brief Readiness calculation

The readiness indicator measures how well the role requirements have been defined. It does not score the hiring manager and is not a validated prediction of hiring success. The language model classifies each area as **confirmed** , **partial** , **conflicting** , or **unknown** ; deterministic n8n code applies multipliers of 1, 0.5, 0.25, and 0 respectively.

Discovery area	Weight	Deterministic safeguard
Role context	10	All nine areas must have a valid status.
Outcomes	15	Fewer than three outcomes cannot be fully confirmed.
Responsibilities and boundaries	10	Fewer than three responsibilities cannot be fully confirmed.
Capability priorities	15	Missing rationales or anchors reduce readiness.
Role-specific evidence; collaboration and support; candidate evidence; working behaviours; logistics and decisions	10 each	Open decisions reduce logistics readiness; any unresolved conflict caps the overall result at 74.

Scores of 90–100 are *Ready for TA review*; 75–89 are *Ready with minor clarifications*; 60–74 indicate *Requirements need refinement*; below 60 is *Not ready to generate*. The result highlights where human clarification is needed and never changes rubric criteria automatically. TA approval remains mandatory.

4.4 Handling vague responses

Manager: “I need someone good with people.”

Assistant: “What must they achieve through other people in this role—for example, resolve customer issues, influence stakeholders, coordinate a campaign, or coach a team? Tell me about one situation they should handle well and what a good result would look like.”

Possible structured result: “Essential: coordinates weekly launch inputs across Brand, Design, and Sales; raises blockers early; confirms owners and deadlines. Evidence: a specific cross-functional example with the candidate’s actions and outcome.”

The transformation uses four steps:

Step	Questioning operation	Stored requirement
Clarify context	With whom, in what situation, and at what frequency?	Stakeholders and work situation
Make observable	What would the person say or do?	Behaviour, not personality label
Tie to outcome	What changes when they do it well?	Business purpose and success evidence
Define evidence	What candidate example or work sample would demonstrate it?	Fair question and scoring anchor

Other required difficult scenarios

- **Vague executive:** “I need a superstar who can do everything.” Ask for the top three outcomes, which work may be deprioritised, and what must be independently deliverable in the first 90 days. If the outcomes require incompatible seniority or disciplines, flag a possible role split for human review.
- **Over-technical entry-level brief: 47 required skills.** Group skills by capability, ask which are needed on day one, require an outcome-based reason for each essential, and move supported learning items to trainable. Preserve alternatives such as “one of Python, JavaScript, or a comparable scripting language.”
- **Culture-first:** “Skills can be taught; I need someone who gets our vibe.” Ask which work behaviour matters, when it appears, and what evidence would distinguish it. Map only an agreed behaviour to a relevant ZURU DNA theme.

5 Capturing technical skills and cultural fit

The system captures cultural fit as **culture contribution through observable work behaviour**. It does not assess whether a person feels socially similar to the team.

Dimension	Technical/craft capability	Culture and working behaviour
Definition	A task the role must perform using relevant knowledge, tools, or judgment	A role-relevant action in a work context linked to an agreed DNA theme
Required fields	Capability, priority, outcome reason, independence, acceptable alternatives, trainability	DNA theme, observable behaviour, situation, expected impact, boundaries
Evidence	Work sample, project, scenario, system explanation, measurable outcome	Specific example of preparation, feedback, constructive challenge, collaboration, or purposeful experimentation
Anchors	Four role-level descriptions from below requirement to exceeds requirement	Four behaviour anchors describing actions, scope, and impact; never personality traits
Prohibited inference	Prestige, years alone, tool keyword count, unsupported equivalence	“Fun”, “like us”, accent, sociability, personality, availability, demographic similarity

5.1 ZURU DNA translation examples

DNA theme	Acceptable observable requirement	Avoid
Shift the Needle	Proposes a small test tied to a campaign goal and identifies the measure that would change the next decision	“Naturally entrepreneurial” or constant availability
Compounding Improvement	Seeks specific feedback, applies it, and records or shares the learning so the same error is not repeated	“Always positive about feedback”
Overprepare & Win	Identifies dependencies, prepares evidence, and simplifies execution before a launch or review	Rewarding unpaid overwork
Radical Candour	Raises a material concern respectfully, with evidence and a constructive alternative	Aggressiveness or extroversion
Collaboration	Coordinates owners, deadlines, hand-offs, and shared credit across relevant teams	“Good culture fit” or friendship similarity
Good Humans Only	Demonstrates expertise, reliable contribution, and respectful work conduct	Judgments about personal worth or protected traits

Every essential technical or behavioural criterion must link to a confirmed role outcome, at least one screening question, and four observable anchors. The validator rejects a rubric if any criterion is not covered by a question.

6 Risk Assessment Matrix

Likelihood and impact are scored from 1 (Low) to 5 (Very high). **Risk score = Likelihood × Impact**. Bands are Low 1–4, Moderate 5–9, High 10–16, and Critical 17–25. Residual scores are recalculated after the listed controls; the arithmetic is shown so the rating can be challenged.

ID	Potential screening bias	Failure mode	Inherent risk	Preventive and detective controls	Residual risk	Human checkpoint
B1	Subjective manager / “culture fit” bias	Vague labels such as “superstar”, “fun”, or “gets our vibe” become hidden screening rules unrelated to job outcomes.	L4 × I5 = 20 Critical	Convert labels into observable work behaviour, context, outcome, and acceptable evidence. Disallow personality and similarity inference. Map DNA only to an agreed behaviour.	L2 × I5 = 10 High	H1 manager confirms the brief; H2 TA challenges unsupported essentials and culture gates.
B2	Prestige and credential proxy bias	School, employer brand, title, years, or a narrow qualification substitutes for demonstrated capability and proxies access.	L3 × I4 = 12 High	Require an outcome-based reason for essentials. Accept equivalent projects, volunteering, portfolios, and adjacent experience. Keep preferred criteria from becoming rejection gates.	L2 × I4 = 8 Moderate	H2 TA approves each essential and its acceptable alternatives before publication.
B3	Language, accent, and communication-style bias	Eloquence, grammar, idiom, accent, tone, or response length is mistaken for competence when it is not job-essential.	L4 × I4 = 16 High	Assess semantic evidence against role anchors. Prohibit inference from accent, tone, emotion, or style. Use accessible wording and score communication only when linked to an outcome.	L2 × I4 = 8 Moderate	H4 reviewer verifies that the quoted evidence and communication criterion are role-relevant.
B4	Unequal evidence and transcription bias	Poor audio, wrong attribution, incomplete notes, inaccessible formats, or uneven interviewer probing creates unequal evidence quality.	L3 × I5 = 15 High	Accept only evidence the uploader is authorised to use; record source and attribution; enforce provenance-based confidence ceilings; require exact excerpts; use null rather than score 1 for missing evidence; permit structured alternatives.	L2 × I4 = 8 Moderate	H3 reviewer verifies authorisation, attribution, and material errors. Blocking or unscoring weak evidence reduces residual impact from 5 to 4.
B5	Automation and confirmation bias	A fluent AI rationale anchors the reviewer; model inconsistency or an early score receives more weight than the underlying evidence.	L4 × I5 = 20 Critical	No aggregate score or hire/reject verdict. Show anchor, quote, confidence basis, and deterministic review priority. Discard unsolicited verdicts; audit corrections and overrides.	L2 × I5 = 10 High	H4 named reviewer owns progression; H5 governance audit samples decisions and monitors overrides and drift.

Residual-risk decision: B1 and B5 remain High because human judgment and model influence cannot be eliminated; they require ongoing approval and audit rather than a claim of full mitigation. B2–B4 reduce to Moderate when the stated controls operate.

6.1 Recommended human checkpoints

Gate	Owner and timing	Required review	System behaviour if unresolved
H1	Hiring manager, before generation	Outcomes, scope, priorities, trainable skills, behaviours, logistics, unresolved decisions	Block generation for missing title, outcomes/scope, essentials, or location/arrangement
H2	TA, before publication/evaluation	Readiness clarification areas, fairness and necessity of essentials, rubric anchors, screening questions, DNA mapping, legal/company wording	Rubric remains draft; evaluator refuses unapproved rubric
H3	Interviewer/reviewer, before evaluation	Candidate name, rubric version, evidence source, uploader authorisation, speaker attribution, obvious transcript errors	Reject inconsistent source/file combinations or transcripts with no matching candidate-labelled turns
H4	TA/hiring panel, after AI report	Exact quotes, applied anchor, confidence basis, concerns, contradictions, and reviewer actions	Report stays “awaiting human review”; no progression output exists
H5	TA governance after the pilot and at scheduled intervals	Reviewer disagreements, corrected AI assessments, inconsistent outcomes, complaints, and changes after model or workflow updates	Review or pause the affected criterion or workflow when repeated errors or unexplained patterns appear

7 Process Flow Diagrams

7.1 Current hiring workflow and pain points




Figure 2: Current-state pain points.

7.2 Proposed AI-augmented workflow




Figure 3: Proposed AI-augmented workflow. Red blocks indicate A.I assisted meanwhile Green blocks represent human checkpoints

8 Candidate evaluation logic

8.1 Score, confidence, and review priority are separate

Output	Meaning	Rule
Criterion score	Match between supplied evidence and the approved 1–4 anchor	AI selects one anchor and must quote exact supporting evidence; missing evidence is null, never 1
Evidence confidence	Strength and reliability of evidence for that criterion	High: direct, specific, complete, consistent; Medium: relevant with a material gap; Low: vague, indirect, contradictory, or missing; source provenance sets a ceiling
Human-review priority	Where a person should focus attention	High: missing/low-confidence evidence or anchor 1 on essential; Medium: anchor 2 on essential or any grounded concern; Low: clear evidence with no concern

A verified attributed transcript allows a high confidence ceiling. Structured notes, pasted answers, or partially verified attribution are capped at medium. Unattributed evidence is capped at low. This is an evidence-strength category, not a probability that the candidate will succeed or that a claim is true.

8.2 Post-interview reviewer actions

Because the evaluator runs after the interview, it does not generate a follow-up question for every criterion. Clear evidence shows no additional action. Missing, contradictory, low-confidence, or attribution-sensitive evidence produces a concrete reviewer action, such as checking the transcript or interviewer notes. Candidate clarification is suggested only at a later authorised stage when the unresolved point is material.

9 Operational controls, rollout, and metrics

9.1 Failed approaches and lessons learned

Early integration experiments were deliberately broader than the final POC. They exposed where production-channel setup would consume effort without improving the core demonstration.

Approach tested	What made it unsuitable for the POC	Lesson and resulting decision
WhatsApp via Meta	Meta app configuration, webhook ownership, public callback requirements, test/production modes, and business-account setup created substantial platform work before the hiring flow could be tested.	Prove the conversation and routing logic through Telegram, which offers a faster bot setup. Retain WhatsApp as a production channel that can reuse the same channel-independent workflows later.
Google Docs output	Google Cloud OAuth consent, redirect-URI management, tester restrictions, and changing temporary tunnel URLs made a simple editable-output feature dependent on external cloud configuration.	Generate reliable local documents for the POC. Treat Google Docs as a later authenticated export option after identity, stable hosting, and organisational Google Cloud approval are available.

These trials reinforced three design choices: validate the hiring logic before adding channel complexity; isolate messaging and document exporters behind replaceable adapters; and distinguish a credible production extension from functionality that must be demonstrated in the take-home prototype.

9.2 Security and governance controls

- Keep model keys and Telegram tokens in n8n credentials; never embed them in workflow exports.

- Treat manager text, uploaded documents, and transcripts as untrusted data, not instructions.
- Minimise candidate data in the POC; use synthetic evidence for demonstration.
- In production, add SSO/RBAC, encryption, retention/deletion policy, audit history, regional/legal review, and vendor data-processing approval.
- Version prompts, model, rubric, and renderer; retain the rubric snapshot used for each report.
- Treat brief readiness as a workflow aid: disclose its deterministic method, preserve the underlying area statuses, and require TA review regardless of score.
- Monitor malformed outputs, model latency/cost, reviewer overrides, and group-level fairness indicators where lawful and appropriate.

9.3 Phased rollout

Phase	Scope	Exit evidence
POC	Synthetic Marketing Intern plus technical and creative scenarios; local n8n; Telegram; manual transcript upload	Complete pack, approved rubric handoff, grounded candidate report, workflow tests and live demo
Controlled pilot	Small set of real vacancies with trained TA reviewers; shadow decisions only	Usability, time saved, revision rate, agreement, override, incident, and fairness review
Production hardening	Persistent state, identity, secure storage, Graph/ATS integration, audit and retention	Security/privacy approval, failure recovery, access testing, monitoring, documented ownership
Scaled operation	Additional teams/regions and controlled model updates	Periodic bias audit, drift review, KPI improvement, rollback procedure, reviewer training

9.4 POC validation evidence and pilot test plan

The POC has been tested for deterministic behaviour and demonstrability, not hiring effectiveness. On 21 September 2026, the final two-workflow repository test suite completed **27/27 checks with 0 failures**. The saved demonstration set also contains both required hiring-pack PDFs, one approved rubric with five criteria and six screening questions, and a 32-cue synthetic interview transcript. These are concrete functional checks; they do not establish model accuracy, fairness, time saved, or better hiring outcomes.

Measure	Current evidence or pilot target	How it is tested
Deterministic workflow validation	Actual: 27/27 automated checks pass; 0 fail	Run both workflow validation test files; covers schemas, required question areas, employment rules, readiness calculation, safe rendering, routing, approval gates, evidence quotes, scores, confidence ceilings, and PDF handling.
Hiring-brief readiness	Actual: weighted calculation, conflict cap, internship normalisation, and permanent-role date controls pass automated tests	Re-run fixed complete, vague, conflicting, internship, and permanent full-time fixtures; verify the same inputs always produce the same score and clarification areas.
Required hiring-pack output	Actual fixture: 2/2 PDFs generated	Confirm a valid job-description PDF and hiring-guide PDF are produced from one approved brief and retain the same rubric ID.
Evaluator demonstration input	Actual fixture: 1 approved five-criterion rubric, 6 questions, and 1 synthetic VTT transcript with 32 cues	Run the post-interview evaluator; verify every criterion is represented once, cited text exists in the source, and the report remains awaiting human review.

Measure	Current evidence or pilot target	How it is tested
End-to-end completion rate	Pilot target: at least 9 of 10 sessions complete without technical assistance	Observe ten hiring managers completing discovery, reviewing the brief, generating the pack, and opening both documents. Record the failed step and reason.
Time to approved hiring pack	Pilot target: median at most 20 minutes	Timestamp first bot interaction and TA approval for each pilot role; report median and range.
Evidence-grounding accuracy	Pilot target: 100% of cited excerpts found verbatim; at least 90% of criterion mappings accepted by two human reviewers	Audit all excerpts automatically, then have two reviewers independently accept or correct each criterion mapping. Report denominator, agreement, and corrections.
Safety gates	Pilot target: 100% of unapproved rubrics, inconsistent source declarations, and unmatched transcript speakers blocked	Run a fixed negative test set before each release; any false acceptance blocks release.

Broader outcome measures such as reduced JD rework, fewer basic-requirement misses, reviewer override patterns, and group-level fairness require a controlled pilot with real baseline data. They should be measured later and must not be presented as results of this take-home POC.

10 Traceability to assignment deliverables

Required deliverable	Location in this document
Information Extraction Framework	Section 4: nine questions, coverage rules, difficult scenarios
Handling vague responses	Section 4.4: complete “good with people” transformation
Technical skills and cultural fit	Section 5: common data model and DNA behaviour examples
Five potential automated-screening biases	Section 6: B1–B5 risk matrix
Mitigation strategies	Section 6: control strategy for every bias
Human checkpoint recommendations	Section 6.1: H1–H5 owners, reviews, and blocking behaviour
Current workflow pain points	Section 7.1: baseline flow and supplied metrics
Proposed AI-augmented workflow	Section 7.2: end-to-end flow with human checkpoints
System architecture	Sections 2–3 and 9: components, boundaries, controls, rollout

AI structures requirements and evidence.
Humans approve the rubric, interpret context, and own every hiring decision.